

# Binoculars Project Agent Guide

## 1) Project Purpose

This repository implements a local, likelihood-based AI text forensics tool inspired by Binoculars.

Primary goal:

- Score markdown text using two related llama.cpp models (observer + performer) and compute:
  - `logPPL` for observer
  - `logXPPL` cross-entropy term
  - $B = \logPPL / \logXPPL$  (Binoculars ratio)

Why this exists:

- `initial-scoping.md` documents that API-only approaches (OpenAI/Ollama top-N logprobs) are approximate and fragile.
- The chosen direction is faithful local scoring with full logits via `llama-cpp-python`.

## 2) Current Project State

Status:

- Functional CLI + GUI prototype with regression coverage for the `v1.1.x` line.
- Core scoring is stable; calibration/classification is not implemented.
- GUI supports iterative rewrite workflows with approximate impact scoring and explicit full re-analysis.
- GUI also supports fast synonym-assisted edits, one-level undo for tracked mutations, and transient status messaging that restores metrics automatically.
- CLI also supports local HTTP API mode via `--api [PORT]` for text-segment scoring (`GET /health`, `POST /score`).
- A dedicated API demo harness is available (`api_demo_harness.py`) with an auto-launch script (`run_api_demo_harness.sh`) for experimentation.
- VS Code extension is active and featureful:
  - chunk-aware `Analyze Chunk` / `Analyze Next Chunk`
  - rewrite selection/line with ranked options
  - colorization + gutter bars + hover diagnostics
  - `Toggle Colorization` runtime command
  - prior contributor faint backgrounds (major contributors only)
  - sidecar state restore (`<doc>.json`)

Latest known commit at time of this guide update:

- 225b75e

### 3) Repository Map

Key files:

- `binoculars.py` : main scoring CLI + GUI application.
- `binoculars.sh` : venv-activating wrapper runnable from any directory.
- `binocular.sh` : alias wrapper (`exec binoculars.sh`).
- `config.binoculars.json` : master profile map ( `fast` / `long` ) + default profile + optional per-profile `max_tokens` override.
- `config.binoculars.llm.json` : optional OpenAI-compatible rewrite backend config for GUI rewrites.
- `config.llama31.cuda12gb.fast.json` : default profile ( `text.max_tokens=4096` ).
- `config.llama31.cuda12gb.long.json` : long profile ( `text.max_tokens=12288` ).
- `README.md` : project overview and CLI/GUI reference.
- `USERGUIDE-GUI.md` : GUI-specific interactive workflow guide.
- `USERGUIDE-VC.md` : VS Code extension workflow guide.
- `USERGUIDE-API.md` : API server + harness walkthrough and request examples.
- `api_demo_harness.py` : Tkinter API demo client for `/health` and `/score`.
- `run_api_demo_harness.sh` : helper script that starts API (if needed) and opens the demo harness.
- `refresh-binoculars-vscode.sh` : local extension refresh helper (compile/package/install + extension-host and daemon restart).
- `vscode-extension/src/extension.ts` : extension UI/decoration/command logic.
- `vscode-extension/src/backendClient.ts` : persistent JSON bridge client.
- `vscode-extension/python/binoculars_bridge.py` : bridge backend process adapter.
- `vscode-extension/package.json` : command/menu/settings manifest.
- `tests/test_regression_v1_1_x.py` : regression checks.
- `tests/fixtures/Athens.md` : stable fixture copy for regression tests.
- `initial-scoping.md` : technical scoping and tuning lessons.

Local assets (present on this machine):

- `models/Meta-Llama-3.1-8B-Q5_K_M-GGUF/meta-llama-3.1-8b-q5_k_m.gguf` (~5.4G)
- `models/Meta-Llama-3.1-8B-Instruct-Q5_K_M-GGUF/meta-llama-3.1-8b-instruct-q5_k_m.gguf` (~5.4G)

### 4) How the Implementation Works

High-level scoring flow in `binoculars.py` :

1. Load JSON config and validate required sections ( `observer` , `performer` ).
2. Read input markdown from file or stdin.
3. Tokenize text with each model in `vocab_only=True` mode.
4. Enforce exact tokenization match (hard fail if mismatch).

5. Infer `n_ctx` (auto when configured as `0`).
6. Load observer with `logits_all=True`, run `eval(tokens)`.
7. Compute observer `logPPL`, write observer logits to memmap.
8. Unload observer (VRAM reduction).
9. Load performer with `logits_all=True`, run `eval(tokens)`.
10. Compute performer `logPPL` (informational) and cross `logXPPL` using observer memmap + performer logits.
11. Compute `B = logPPL(observer) / logXPPL(observer, performer)`.
12. Emit text or JSON output.
13. Remove cache unless `cache.keep=true`.

Design choice:

- Models are loaded sequentially (never concurrently) to reduce VRAM pressure.

## 5) Config Profiles and Intent

`config.llama31.cuda12gb.fast.json`:

- `max_tokens: 4096`
- `offload_kqv: true`
- `n_batch: 1024`

`config.llama31.cuda12gb.long.json`:

- `max_tokens: 12288`
- `offload_kqv: true`
- `n_batch: 1024`

Both shipped profiles:

- Use Llama 3.1 8B base + instruct Q5\_K\_M sibling models.
- Use `n_ctx: 0` (auto = token count).
- Use cache dtype `float16`.

Master config override behavior:

- `config.binoculars.json` may define each profile as:
  - string path, or
  - object with:
    - `path` : concrete config JSON
    - `max_tokens` : optional non-negative override for `text.max_tokens`

Optional rewrite backend config (`config.binoculars.llm.json`):

- If file missing or disabled: use internal performer rewrite generation.
- If present and reachable: use configured OpenAI-compatible endpoint.
- If present but invalid/unreachable at runtime: auto-fallback to internal generation.

## 6) Environment and Bootstrap

Observed local state:

- Python in repo venv: 3.10.12 .

Baseline dependencies:

```
venv/bin/pip install numpy llama-cpp-python
```

Optional dependency for richer synonym lookup:

```
venv/bin/pip install nltk
venv/bin/python - <<'PY'
import nltk
nltk.download('wordnet', quiet=True)
nltk.download('omw-1.4', quiet=True)
PY
```

Example scoring run:

```
venv/bin/python binoculars.py --config fast your_doc.md --json
```

GUI run:

```
venv/bin/python binoculars.py --config fast --gui your_doc.md
```

API server run:

```
venv/bin/python binoculars.py --config fast --api 8765
```

API harness launcher:

```
./run_api_demo_harness.sh
```

## 7) Output Contract

JSON output includes:

- `input` metadata (chars, tokens, transitions)
- `observer` ( logPPL , PPL )
- `performer` ( logPPL , PPL )
- `cross` ( logXPPL , XPPL )

- `binoculars.score ( B )`
- `cache details`

Important:

- Script returns scores only. No built-in threshold classifier labels.
- GUI focuses on iterative editing + rescoreing:
  - `Analyze` computes exact metrics for the active analyzed chunk and refreshes heatmap coverage.
  - rewrites show approximate impact; exact B requires Analyze.
  - status explicitly marks B stale after edits/rewrites.
  - non-analysis status messages are transient and then restore metrics (Undo success is intentionally brief).

GUI rewrite behavior:

- Right-click on red (`LOW`) scored segments to request 3 rewrites.
- Or highlight a block (multi-line) and right-click for block rewrites.
- Selection rewrites:
  - are rounded to full lines
  - are clamped to scored/analyzed text if needed
  - preserve unchanged source lines when model output accidentally omits/collapses them
- Popup supports scrolling and keyboard selection (`1 / 2 / 3`, `Q / Esc` to cancel).
- Options are sorted by expected B increase (more human-like first).

GUI synonym behavior:

- Left-click a word in the left pane to trigger synonym lookup after a short debounce.
- Synonym panel (bottom of right pane) shows up to 9 options in 3 columns with buttons `1 .. 9`.
- Source order: local fallback -> WordNet (if installed) -> Datamuse API fallback.
- Applying a synonym is tracked as one undoable mutation.

GUI undo behavior (single level):

- Toolbar has `Undo` button.
- Supported tracked operations:
  - selected-block delete (`Delete` or `Backspace`),
  - synonym replacement,
  - red-segment rewrite replacement,
  - highlighted-block rewrite replacement.
- Undo is invalidated if document text changes after the tracked operation.

GUI identity details:

- Window/app name is set to `Binoculars` (including Linux WM class/appname hints).
- GUI icon is drawn in code (owl with large eyes) via Tk `PhotoImage`.

API mode behavior:

- `--api [PORT]` runs a local HTTP server bound to `127.0.0.1`.
- Health endpoints: `GET /`, `GET /health`, `GET /healthz`.
- Scoring endpoint: `POST /score` with required JSON field `text`.
- Optional `/score` request fields:
  - `input_label` (string),
  - `diagnose_paragraphs` (boolean),
  - `diagnose_top_k` (integer),
  - `need_paragraph_profile` (boolean).
- `/score` response returns:
  - `{ "ok": true, "result": <standard JSON scoring object> }`,
  - plus `paragraph_profile` when requested.

VS Code extension behavior (current):

- Commands:
  - `Binoculars: Enable`
  - `Binoculars: Disable`
  - `Binoculars: Analyze Chunk`
  - `Binoculars: Analyze Next Chunk`
  - `Binoculars: Analyze All`
  - `Binoculars: Rewrite Selection`
  - `Binoculars: Clear Priors`
  - `Binoculars: Toggle Colorization`
  - `Binoculars: Restart Backend`
- Color model:
  - Major LOW/HIGH contributors are colorized.
  - Minor contributors are rendered neutral (`light gray` in dark theme, `black` in light theme).
  - Gutter bars remain available independently of text colorization.
- Prior overlays:
  - After re-analysis, prior contributor overlays are captured as faint backgrounds.
  - Prior overlays are restricted to prior major contributors (`top-k` LOW/HIGH), not minor rows.
  - `Clear Priors` clears prior overlays without deleting current analysis state.
- Toggle colorization:
  - `Toggle Colorization` hides/shows text overlays at runtime.
  - Re-enabling restores overlays from in-memory state (including prior/edited backgrounds when present).
- Hover behavior:
  - Contributor/stale hovers use a shared delayed reveal gate for major and minor rows (currently ~1.15s).
  - Same-segment suppression reduces immediate re-pop in the same segment.
- Persistence:
  - Sidecar save/load includes chunk state, edited ranges, rewrite ranges, and prior overlay ranges.

- `priorChunkB` remains in-session and is intentionally not restored from persisted state.
- Estimate behaviors:
  - Rewrite popup options show approximate `approx_B` / `delta_B` using local observer logPPL rescore + transition-normalized projection while holding baseline cross term fixed.
  - Manual typing in stale analyzed text triggers debounced live forecast (`Est. B`) in status using observer-only chunk-start rescore with baseline cross term held fixed.
  - Daemon keeps an observer model warm for live-estimate responsiveness and unloads it automatically after idle timeout.
  - Both estimate paths are directional guidance; explicit `Analyze` / `Analyze Next` / `Analyze All` remain authoritative for exact checkpoint metrics.

## 8) Lessons Learned / Gotchas

1. Full-logit requirement is non-negotiable for faithful Binoculars.
  - API top-k logprob approaches are approximate and biased.
2. Tokenizer/vocab alignment is critical.
  - Script hard-fails on observer/performer tokenization mismatch.
3. Memory pressure is dominated by `(tokens * vocab)` with `logits_all=True`.
  - `text.max_tokens` is the primary safety valve.
4. `n_ctx: 0` auto-sizing avoids over-allocation and many avoidable failures.
5. GUI rewrite scoring is intentionally approximate between analyses.
  - Approximate option ranking is local observer-logPPL based.
  - Cross term is not recomputed until Analyze.
6. Markdown is treated as text.
  - Rendering is for convenience; scoring is raw text-token based.

## 9) Chunk-Aware GUI Analysis (Implemented Behaviour)

Chunk-aware large-file analysis is implemented in the GUI.

### 9.1) Analyze and Analyze Next Semantics

1. First `Analyze`:
  - Starts at document char `0`.
  - Scores forward until token/memory limit for that run.
2. `Analyze Next`:

- Starts at contiguous covered tail (`analysis_covered_until`).
- Scores the next token-limited chunk.
- Remains available until contiguous coverage reaches end-of-document.

3. Later `Analyze` runs:

- Resolve active chunk.
- Start scoring at `active_chunk.char_start` (not cursor char).
- Recompute that chunk from its start boundary.

## 9.2) Active Chunk Resolver Priority

Resolver order is deterministic:

1. Current selection overlap with analyzed chunks (largest overlap wins).
2. Else chunk containing visible insert/cursor line.
3. Else chunk with largest overlap against visible line window.
4. Else nearest analyzed chunk by char distance.

Status metrics and rewrite approximation baselines use this active chunk.

## 9.3) Chunk Boundary Mutability (Important)

Chunks are not immutable “mini-documents”.

- Chunk metrics are stored per chunk descriptor.
- Coverage intervals are merged for rendering and unscored complements.
- On overlap, old descriptor(s) are replaced by the newest re-analysis descriptor.
- Therefore chunk end boundaries may move after edits.

Example:

- First pass may produce chunk 1 covering lines `1-999`.
- User at line `999` presses `Analyze`.
- Analyzer restarts at chunk-1 start (line 1), not line 999.
- After edits, same run may now end at line `972` or `1031` due to token-density changes.

## 9.4) Rewrite Approximation in Chunk Context

- Rewrite requests target selected span/segment as usual.
- Baseline metrics are resolved from the request/active chunk.
- Approximate scoring context is clamped to chunk bounds when available.

## 9.5) Rendering Model

- Combined annotations from all chunk profiles are rendered globally.
- Unscored intervals are computed as complement of merged scored coverage.
- Left gutter bars and preview backgrounds reflect multi-chunk state.

## 10) Known Gaps / Next Development Priorities

Priority backlog:

1. Add calibration pipeline:
  - dataset runner + threshold selection + FPR/TPR reporting.
2. Add sliding-window scoring:
  - support very long docs without full-doc logits materialization.
3. Add tests:
  - synthetic math checks for perplexity/cross-perplexity and rewrite post-processing guards.
4. Add dependency pinning:
  - `requirements.txt` or `pyproject.toml`.
5. Add reproducible benchmark script:
  - throughput/memory across profiles and input lengths.

## 11) Agent Operating Notes

When resuming work:

1. Verify environment first:
  - dependencies, model paths, writable cache directory.
2. Preserve core math semantics unless intentionally changing them.
  - Any changes to token alignment, cross-entropy math, or truncation behavior must be explicit and documented.
3. Preserve sequential model loading unless redesign is intentional and benchmarked.
4. For GUI rewrite changes, keep user control explicit.
  - Do not auto-run full Analyze after each rewrite.
  - Keep approximate impact messaging clear.
5. If classification labels are added in future:
  - keep raw numeric outputs and expose calibration metadata.

## 12) Non-Goals (Current)

Not currently in scope:

- Definitive authorship claims.
- Remote API-only detector approximations for scoring core.
- Production web service deployment.

## 13) VS Code Marketplace Roadmap

Goal:

- Publish `vscode-extension` to the VS Code Marketplace in a way that non-technical users can install and run it successfully.

### 13.1) Publishing Prerequisites

1. Publisher setup:

- Create/verify Azure DevOps publisher identity for Marketplace.
- Generate Personal Access Token (PAT) with Marketplace publish permissions.
- Add secure local/CI publishing flow (`vsce publish` or GitHub Action).

2. Manifest hardening (`vscode-extension/package.json`):

- Add `repository` field (currently missing warning during packaging).
- Keep `displayName`, `description`, `categories`, icon, and command titles user-friendly.
- Ensure license path is valid for packaged extension.

3. Release hygiene:

- Add `CHANGELOG.md`.
- Define semantic versioning and release notes process.
- Add automated prepublish checks (`npm run compile`, smoke checks).

### 13.2) Installation Friction Assessment (Config Streamlining)

Short answer:

- Yes, streamlining is likely needed for broad adoption.
- For expert users, current explicit paths can work; for general users, manual path/model setup is too error-prone.

Why:

- Current defaults point to machine-specific absolute paths.
- Users must have Python, dependencies, config JSON, and GGUF model paths aligned.
- Without guided setup, first-run failure rate will be high outside the current dev environment.

### 13.3) Accessibility Roadmap (Broader Audience)

Phase 1 (minimum for Marketplace launch):

1. Replace machine-specific defaults with portable defaults:
  - Use  `${workspaceFolder}` where appropriate.
  - Leave optional paths empty when unknown and detect at runtime.

2. Add first-run preflight + guided setup command:

- Validate Python executable, bridge script, config file, model files.
- Show actionable one-click fixes/open-settings shortcuts.
- Persist discovered valid paths automatically.

3. Improve error UX:

- Human-readable diagnostics in notifications and output channel.
- Clear distinction between missing dependency, missing model, bad config, and backend startup failure.

Phase 2 (recommended post-launch):

1. Setup wizard:

- Multi-step onboarding UI for selecting profile/config/models.
- “Test backend” button before first analyze.

2. Optional quickstart profile:

- Ship a template config and docs that minimize manual edits.
- Provide explicit “local-only” and “external rewrite backend” setup paths.

3. Telemetry-free health metrics (local only):

- Count setup failures in session and surface targeted help (no remote telemetry required).

Phase 3 (best usability):

1. Dependency/bootstrap helper:

- Command to create/check venv and install Python dependencies.
- Optional model path validator/downloader integration (if licensing/distribution allows).

2. Cross-platform QA matrix:

- Linux/macOS/Windows first-run validation with clean machines.

### 13.4) Acceptance Criteria Before Public Marketplace Push

1. Fresh-machine install succeeds with guided steps (no code edits required).
2. User can run first `Analyze Chunk` within a short onboarding flow.
3. Common failures provide direct remediation links/commands.
4. Documentation is aligned:

- Root README.md
- USERGUIDE-API.md
- vscode-extension/README.md
- USERGUIDE-VC.md